

KESTRELATTICE

Governed Agent Mesh | Micro-Product

Incident Communication Template Pack

What to tell the team, the customer, and the log during an agent incident. A complete, zero-placeholder draft for teams governing AI agents in production: policy gates, audit trails, cost controls, permissioning, and incident response, written so you can adopt it as-is.

A complete, ready-to-use governance document. Adopt it as-is, customize the bracketed fields to your environment, and put it to work today.

Instant delivery | Yours to keep, forever.

Incident Communication Template Pack

For AI Agent Governance in Production Environments **Version 1.0 License:** MIT (Modifiable for internal use)

1. Executive Summary

This template pack standardizes incident communication for AI agents in production, ensuring **compliance, traceability, and operational rigor**. It includes: - **Incident Report Template** (initial triage) - **Escalation Protocol** (multi-tiered governance) - **Post-Incident Review (PIR) Template** (root cause + remediation) - **Cost & Resource Impact Log** (financial governance) - **Permission & Audit Log** (access control)

Each template is **field-populated** with real-world governance controls (e.g., policy gates, SLAs, audit trails).

2. Incident Report Template

Purpose: Capture, classify, and triage AI agent incidents in real time.

Field	Definition	Example Value
Incident ID	Auto-generated UUID (e.g., INC-2024-0001)	INC-2024-0042
Timestamp	UTC timestamp of detection (ISO 8601)	2024-05-15T14:32:17Z
Agent Name	Fully qualified AI agent identifier (e.g., prod/finance/credit-scoring-v2)	prod/hr/compliance/audit-logger-v1
Incident Type	Predefined severity + category (see Appendix A)	Severity: CRITICAL Category: Data Leakage
Description	Step-by-step reproduction (include logs, API calls, and user inputs)	"Agent prod/hr/compliance/audit-logger-v1 exposed PII in error logs via API endpoint GET /v1/debug/logs to IP 192.168.1.100 (internal subnet)."
Root Cause Hypothesis	Initial guess (updated in PIR)	"Missing X-Request-Auth header validation in /debug endpoint."
Impact Assessment	Data: # records exposed; Cost: \$; Uptime: %; Reputation: N/A	Data: 12,450 employee records (SSN, salary); Cost: \$45,000 (forensic audit); Uptime: 0% (service halted)
Detected By	Role (e.g., "Monitoring Engineer," "User Report")	"Monitoring Engineer"
Detected At	Time of discovery	2024-05-15T14:45:00Z
Initial Mitigation	Actions taken (e.g., "Disabled endpoint," "Rotated API keys")	"Endpoint GET /v1/debug/logs hard-disabled via Kubernetes patch; keys rotated via Vault."
Owner	Assigned team (e.g., "SRE," "Compliance")	"SRE (On-Call)"

Field	Definition	Example Value
Escalation Path	Tier 1 -> Tier 2 -> Tier 3 (see Section 4)	"Tier 1: SRE Lead -> Tier 2: Security -> Tier 3: Legal"
Attachments	Logs, screenshots, or relevant artifacts (hashed links)	s3://prod-logs/2024-05-15/incident-42/logs.zip (SHA256: a1b2c3...)
SLA Compliance	Was response time within SLA? (e.g., "CRITICAL: 15 mins")	"No (detected at 14:32, SLA breach at 14:47)"

Appendix A: Incident Types | Severity | Category | Definition |

Severity	Category	Definition
CRITICAL	Data Leakage	Unauthorized exposure of PII/PCI/confidential data.
CRITICAL	System Unavailability	Agent fails for >50% of users or critical workflows.
HIGH	Policy Violation	Agent outputs violate compliance rules (e.g., GDPR, HIPAA).
HIGH	Cost Spiral	Unauthorized API calls exceed budget by >20%.
MEDIUM	Performance Degradation	Latency >2x baseline or throughput <80%.
LOW	Non-Critical Bug	Minor functionality issue (e.g., "Agent returns 404 for valid input").

3. Escalation Protocol

Purpose: Define clear governance gates for AI agent incidents.

Tier 1: Immediate Response (SRE/DevOps)

- **Trigger:** All CRITICAL incidents.
- **Actions:**
 - Disable affected agent endpoints.
 - Rotate all credentials (API keys, service accounts).
 - Notify **on-call engineer** via PagerDuty.
- **SLA:** Mitigate within **15 minutes**.

Tier 2: Security Review (SOC/Compliance)

- **Trigger:** Data leakage, policy violations, or cost spikes >\$10K.
- **Actions:**
 - Conduct forensic analysis (retention: **30 days**).
 - Escalate to legal if PII exposed.
 - Update **Incident Report** with findings.
- **SLA:** Review within **2 hours**.

Tier 3: Executive Notification (CISO/VP Engineering)

- **Trigger:** CRITICAL incidents with **reputational risk** (e.g., media exposure).
- **Actions:**
 - Draft internal communication (template below).
 - Coordinate with PR team.
 - Mandate **Post-Incident Review (PIR)**.
- **SLA:** Notification within **30 minutes**.

Internal Communication Template (Tier 3) Subject: Urgent: AI Agent Incident - [Incident ID] - [Severity] **Body:**

Team,

A **CRITICAL incident** (ID: **INC-2024-0042**) has been detected in **Agent prod/hr/compliance/audit-logger-v1**. **Impact:** 12,450 employee records exposed via debug endpoint.

Current Status: Mitigated (endpoint disabled; forensic audit in progress). **Next Steps:** - [] Legal review (target: 2024-05-16 09:00 UTC). - [] PIR completion by 2024-05-18. - [] All teams: Audit logs for unauthorized access.

Do NOT discuss externally. Questions? Escalate to **security@company.com**.

--[Your Name], VP Engineering

4. Post-Incident Review (PIR) Template

Purpose: Document root cause, remediation, and process improvements.

Field	Definition	Example Value
Incident ID	Reference to initial report	INC-2024-0042
Review Date	UTC timestamp of completion	2024-05-18T10:15:00Z
Root Cause	Final analysis (supported by evidence)	"Missing X-Request-Auth header validation in /debug endpoint due to merge conflict in PR #1234 (unreviewed)."
Immediate Fix	Permanent remediation (e.g., "Added RBAC to endpoint")	"Deployed Kubernetes NetworkPolicy to block /debug from non-admin IPs; enabled audit logging for all API calls."
Long-Term Fix	Process change (e.g., "Enforce PR reviews for debug endpoints")	"Added debug-endpoint label to require 2+ approvals in GitHub; automated header validation tests."
Cost Impact	Total financial loss (forensic, downtime, fines)	\$72,000 (forensic audit + \$25K regulatory fine for GDPR violation)
Time Impact	Agent downtime + recovery time	4 hours (mitigation) + 12 hours (forensic)
Process Gaps	What failed? (e.g., "Missing pre-deploy security scan")	"No automated check for debug endpoint exposure in CI pipeline."
Actions Taken	Corrective measures (assign owners + deadlines)	- Owner: DevSecOps; Deadline: 2024-05-20; Action: "Add debug-endpoint scanner to CI."
> - Owner: Compliance; Deadline: 2024-05-22; Action: "Train engineers on PII handling."		
Audit Trail	Links to relevant logs/emails	s3://audit/2024-05/incident-42/forensic-report.pdf
Approved By	Role (e.g., "Security Lead," "VP Engineering")	"Approved by: Sarah Chen, CISO"

5. Cost & Resource Impact Log

Purpose: Track financial and operational costs tied to incidents.

Incident ID	Agent	Category	Cost Driver	Amount (USD)	Notes
INC-2024-0042	prod/hr/audit-logger-v1	Data Leakage	Forensic audit	\$45,000	Engaged [Firm] for PII breach analysis.
			Regulatory fine (GDPR)	\$25,000	Penalty for inadequate access controls.
			Downtime (engineering)	\$12,000	4 hours of on-call labor.
			PR damage control	\$5,000	External communications team.
Total				\$87,000	

Budget Impact: This exceeds the \$50K quarterly AI incident budget. Escalate to Finance for approval.

6. Permission & Audit Log

Purpose: Track who accessed incident data and when.

| Timestamp (UTC) | User | Role | Action | **Agent

KESTRELATTICE

The governed agent mesh for production AI.

kestrelattice.com