

KESTRELATTICE

Governed Agent Mesh | Micro-Product

New Model Version Approval Flow

Approve a model upgrade like a deployment: gates, rollback criteria, and a named owner. A complete, zero-placeholder draft for teams governing AI agents in production: policy gates, audit trails, cost controls, permissioning, and incident response, written so you can adopt it as-is.

A complete, ready-to-use governance document. Adopt it as-is, customize the bracketed fields to your environment, and put it to work today.

Instant delivery | Yours to keep, forever.

NEW MODEL VERSION APPROVAL FLOW

For AI Agents in Production Environments Version: 1.2 (Governance & Compliance) **Effective Date:** [EFFECTIVE DATE] **Owner:** [GOVERNANCE ADMINISTRATOR] **Last Reviewed:** [LAST REVIEW DATE]

1. PURPOSE

This document establishes a **standardized approval workflow** for deploying new or updated AI model versions into production environments governed by [TEAM NAME]. The flow ensures alignment with **compliance requirements, operational risk controls, and cost governance** while maintaining auditability. It applies to all AI agents, including but not limited to: - Generative AI models (LLMs, embeddings, agents) - Vision/Audio models (computer vision, speech recognition) - Custom fine-tuned or proprietary models - Third-party API-based AI services integrated into [PRODUCT/APPLICATION NAME]

Failure to follow this process may result in **operational disruptions, regulatory penalties, or financial exposure** due to unapproved model behavior.

2. SCOPE

2.1 In Scope

- New model versions (v1.0 -> v1.1, etc.)
- Major updates (e.g., architecture changes, new training data)
- Third-party model integrations (e.g., new LLM provider)
- Model version rollouts to **staging -> production** environments
- **Cost-sensitive** model deployments (e.g., high-token-count generation)

2.2 Out of Scope

- Minor bug fixes (defined as non-functional changes)
- Non-production sandboxing (governed by [SANDBOX POLICY])
- Models used in **research-only** contexts (governed by [RESEARCH POLICY])

3. APPROVAL FLOW DIAGRAM

(Visual representation follows; text workflow below)

3.1 Step 1: Request Submission

Submitter: [REQUESTER NAME] (e.g., Data Scientist, Product Manager) **Form:** [MODEL REQUEST FORM URL] (pre-populated with auto-generated fields)

Field	Requirement	Example Value
Model Name	Exact version identifier (e.g., mistral-7b-v2.1)	gpt-4o-mini-v2
Request Type	New / Update / Third-Party / Cost-Optimization	Update
Environment Target	Staging / Production / [CUSTOM ENVIRONMENT]	Production (Phase 2)

Field	Requirement	Example Value
Requester's Manager Approval	Mandatory digital signature (via [SIGNATURE TOOL])	[MANAGER NAME]'s signature link
Risk Assessment Score*	Auto-calculated from:	
- Compliance Risk (0-100)	Regulatory exposure (e.g., GDPR, HIPAA)	32
- Operational Risk (0-100)	Potential for out-of-distribution failures	18
- Cost Risk (0-100)	Estimated monthly spend increase	45
Total Risk Score	$= (\text{Compliance} \times 0.4) + (\text{Operational} \times 0.35) + (\text{Cost} \times 0.25)$	30.7 (! Requires [TEAM NAME]'s review)

**Risk scoring uses [TEAM NAME]'s proprietary model governed by [RISK ASSESSMENT POLICY].*

3.2 Step 2: Compliance & Security Review

Reviewer: [COMPLIANCE OFFICER] or [SECURITY TEAM LEAD] **Deadline:** [REQUESTER'S MANAGER]'s deadline **minus 48 hours**

Checklist Item	Action Required	Evidence Required
Model Card Compliance	Verify [MODEL CARD URL] includes:	
- Purpose limitation clause	[x]/[]	Screenshot or signed attestation [BIAS REPORT URL] [DATA GOVERNANCE POLICY] reference
- Bias mitigation statement	[x]/[]	
- Data governance section (e.g., PII handling)	[x]/[]	
Third-Party Vendor Due Diligence	If using external models:	
- Signed [VENDOR AGREEMENT] with [TEAM NAME]'s legal team	[x]/[]	[AGREEMENT ID]
- Model's SOC 2 Type II certification (if applicable)	[x]/[]	[CERTIFICATION URL]
Data Protection Impact Assessment (DPIA)	If processing sensitive data:	
- Completed [DPIA FORM] with risk mitigation plan	[x]/[]	[DPIA REFERENCE NUMBER]
Cost Threshold Check	Compare against [TEAM NAME]'s [COST ALERT THRESHOLDS]:	
- Monthly spend > [\$[COST THRESHOLD]] requires [FINANCE TEAM]'s approval	[x]/[]	[FINANCE APPROVAL LINK]

Approval Path: - **Pass:** [COMPLIANCE OFFICER] grants approval via [APPROVAL TOOL]. - **Fail:** Escalate to [TEAM NAME]'s [GOVERNANCE COMMITTEE] with justification.

3.3 Step 3: Technical & Operational Review

Reviewer: [ENGINEERING LEAD] or [SYSTEMS ARCHITECT] **Deadline:** [COMPLIANCE OFFICER]'s deadline minus 24 hours

Checklist Item	Action Required	Evidence Required
Model Performance Baseline	Compare against MODEL	
- Latency (p99 < [LATENCY THRESHOLD] ms)	[x]/[]	[PERFORMANCE REPORT URL]
- Accuracy (confidence interval > [ACCURACY THRESHOLD]%)	[x]/[]	[A/B TEST RESULTS]
Cost Optimization	Verify no cost anomalies:	
- Token usage < [\$[TOKEN COST THRESHOLD]/request]	[x]/[]	[COST ANALYSIS TOOL] output
- Caching enabled for repeat queries	[x]/[]	[CACHE CONFIGURATION URL]
Incident Response Plan	Attach updated [INCIDENT RESPONSE PLAYBOOK] for this model version.	
- Escalation contacts (on-call rotations)	[x]/[]	[PAGERDUTY ALERT RULES]
- Rollback procedure documented	[x]/[]	[ROLLBACK PROCEDURE URL]

Approval Path: - **Pass:** [ENGINEERING LEAD] approves via [APPROVAL TOOL] with **versioned notes**. - **Fail:** Requires [TECH LEAD]'s sign-off or [GOVERNANCE COMMITTEE] review.

3.4 Step 4: Staging Deployment & Monitoring

Owner: [DEVOPS TEAM] **Duration:** [STAGING PERIOD] (e.g., 72 hours)

Requirement	Action	Verification Tool
Canary Release	Deploy to 5% of [STAGING ENVIRONMENT] traffic with [MONITORING TOOL].	[NEW RELIC / PROMETHEUS] alerts
Anomaly Detection	Set alerts for:	
- Error rate > [ERROR RATE THRESHOLD]%	[x]/[]	[ERROR LOG STREAM]
- Latency spikes > [LATENCY ALERT THRESHOLD] ms	[x]/[]	[APM TOOL]
- Cost spikes > [\$[COST ALERT THRESHOLD]]/hour	[x]/[]	[FINANCE DASHBOARD]
Audit Trail	Ensure all API calls logged to [AUDIT LOG STORE] with:	
- Requester ID	[x]/[]	[LOG QUERY]

Requirement	Action	Verification Tool
- Model version	[x]/[]	[LOG SAMPLE]
- Timestamp (UTC)	[x]/[]	[TIMESTAMP EXAMPLE]

Escalation Triggers: - **Critical:** [DEVOPS LEAD] must halt deployment if: - **Hallucination rate** > [HALLUCINATION THRESHOLD]% (defined in [QUALITY POLICY]). - **Regulatory violation** detected (e.g., bias, privacy leak). - **Warning:** [TECH LEAD] reviews if cost exceeds [WARNING THRESHOLD].

3.5 Step 5: Production Approval

Reviewer: [PRODUCT MANAGER] or [HEAD OF AI] **Final Checklist:**

Item	Action	Approval Source
Staging Results	All metrics within [ACCEPTANCE CRITERIA].	[STAGING REPORT]
Cost Approval	If > [\${FINANCE THRESHOLD}], signed by [FINANCE DIRECTOR].	[FINANCE APPROVAL]
Incident Response Plan	Confirmed with [SECURITY TEAM].	[PLAYBOOK SIGN-OFF]
Compliance Sign-Off	[COMPLIANCE OFFICER]'s final approval.	[COMPLIANCE TOOL]
Change Management	Updated [CHANGE LOG] with:	
- Model version	[x]/[]	[CHANGE LOG URL]
- Deployment timestamp	[x]/[]	[TIMESTAMP]
- Approvers' names	[x]/	

KESTRELATTICE

The governed agent mesh for production AI.

kestrelattice.com