

KESTRELATTICE

Governed Agent Mesh | Micro-Product

Observability Scorecard

Tool

Vendor

Score tracing and monitoring tools on the audit evidence they actually produce. A complete, zero-placeholder draft for teams governing AI agents in production: policy gates, audit trails, cost controls, permissioning, and incident response, written so you can adopt it as-is.

A complete, ready-to-use governance document. Adopt it as-is, customize the bracketed fields to your environment, and put it to work today.

Instant delivery | Yours to keep, forever.

OBSERVABILITY TOOL VENDOR SCORECARD

For Buyers Governing AI Agents in Production [EFFECTIVE DATE] [ORGANIZATION NAME]
[BUYER TEAM NAME] [BUYER ROLE] [CONTACT EMAIL] [SCORECARD VERSION]

1. SCORECARD OVERVIEW

This document evaluates observability tools for governing AI agents in production. It assesses compliance with [ORGANIZATIONAL AI GOVERNANCE POLICY] and [REGULATORY FRAMEWORK] (e.g., GDPR, CCPA, NIST AI RMF, or internal risk frameworks). The scorecard covers **policy enforcement, auditability, cost controls, permissioning, and incident response**--critical for AI agent governance.

Purpose: - Compare vendors against [ORGANIZATION'S AI GOVERNANCE PRIORITIES]. - Identify gaps in observability, compliance, and operational resilience. - Guide procurement decisions for tools managing AI agents in production.

Scope: - Applies to tools monitoring, logging, and governing AI agents (e.g., LLM-based agents, autonomous systems, or hybrid workflows). - Excludes vendor-neutral frameworks (e.g., Prometheus) unless integrated into the tool.

2. VENDOR INFORMATION

Field	Value
Vendor Name	[VENDOR NAME]
Tool Name	[TOOL NAME]
Primary Contact	[NAME], [TITLE], [EMAIL]
Pricing Model	[e.g., Per-agent, Per-query, Flat-rate, Usage-based]
Deployment Options	[On-prem, Cloud, Hybrid, Edge]
Integration Capabilities	[APIs, SDKs, Plugin Ecosystem, Custom Connectors]

3. SCORECARD SECTIONS

3.1 POLICY ENFORCEMENT & COMPLIANCE

(Ensures AI agents adhere to organizational rules and external regulations.)

Criteria	Scoring (1-5)	Notes	Evidence Required
Rule Engine Support	[1-5]	Does the tool support policy gates (e.g., input validation, output filtering, hallucination detection)?	[Attach sample ruleset or documentation]

Criteria	Scoring (1-5)	Notes	Evidence Required
Regulatory Alignment	[1-5]	Compliance with [REGULATORY FRAMEWORK] (e.g., bias mitigation, data sovereignty).	[Certifications, audit logs, or compliance reports]
Dynamic Policy Updates	[1-5]	Can policies be updated without redeployment of agents?	[Documentation on live-configuration APIs]
Explainability for Policies	[1-5]	Does the tool provide audit trails for policy violations?	[Sample incident report or policy-execution log]
Third-Party Policy Validation	[1-5]	Supports external policy feeds (e.g., from governance platforms)?	[API specs or integration guides]

3.2 AUDITABILITY & TRACEABILITY

(Ensures all AI agent actions are logged, immutable, and queryable.)

Criteria	Scoring (1-5)	Notes	Evidence Required
Immutable Audit Logs	[1-5]	Logs are time-stamped, tamper-proof, and stored per [AUDIT RETENTION POLICY].	[Sample log export or SLA]
Agent Activity Tracking	[1-5]	Tracks input/output, prompts, context, and agent decisions in real-time.	[Demo of query interface or API response]
Lineage for Data & Decisions	[1-5]	Can trace how an output was derived from inputs?	[Visualization of decision trees or provenance graphs]
Third-Party Audit Access	[1-5]	Supports external auditors (e.g., regulators, internal compliance teams)?	[Access control documentation]
Log Query & Export Capabilities	[1-5]	Supports SQL-like queries or custom filters for logs?	[API docs or query examples]

3.3 COST CONTROLS & BUDGET MANAGEMENT

(Prevents runaway costs from AI agent usage.)

Criteria		Scoring (1-5)	Notes	Evidence Required
Usage-Based Alerts	Cost	[1-5]	Triggers alerts at [COST ALERT THRESHOLD]% of budget.	[Integration with finance tools or custom alerting]
Query/Invocation Limits		[1-5]	Enforces rate limits per agent or user?	[Documentation on throttling or quota management]
Cost Anomaly Detection		[1-5]	Identifies unexpected cost spikes (e.g., rogue agents)?	[Sample anomaly report or ML-based detection]
Budget Ownership & Approvals		[1-5]	Assigns cost budgets to teams with approval workflows?	[Workflow diagrams or integration with ERP]
Cost Transparency Reporting		[1-5]	Provides detailed cost breakdowns by agent, user, or query?	[Sample cost report or API response]

3.4 PERMISSIONING & ACCESS CONTROL

(Ensures only authorized entities interact with AI agents.)

Criteria		Scoring (1-5)	Notes	Evidence Required
Role-Based Control (RBAC)	Access	[1-5]	Supports [ORGANIZATIONAL RBAC MODEL] (e.g., requester, approver, auditor).	[IAM integration docs or policy examples]
Temporary Tokens	Access	[1-5]	Issues time-limited tokens for third-party audits?	[Token management API or workflow]
Agent-Specific Permissions		[1-5]	Granular control over which agents users can invoke?	[Demo of permission assignment interface]
Audit Logs for Access Events		[1-5]	Logs who accessed what and when?	[Sample access log or SLA]
Integration with IAM Systems		[1-5]	Syncs with [IAM SYSTEM] (e.g., Okta, Azure AD, or custom LDAP)?	[SSO documentation or integration guide]

3.5 INCIDENT RESPONSE & REMEDIATION

(Ensures rapid detection and mitigation of AI agent failures.)

Criteria	Scoring (1-5)	Notes	Evidence Required
Real-Time Incident Alerts	[1-5]	Notifies [INCIDENT RESPONSE TEAM] at [ALERT THRESHOLD] severity.	[Integration with PagerDuty/SLACK or alerting docs]
Automated Remediation Actions	[1-5]	Can auto-stop or quarantine misbehaving agents?	[Playbook examples or API for remediation]
Post-Incident Analysis	[1-5]	Provides root cause reports with actionable insights?	[Sample report or ML-based analysis]
Incident Simulation Testing	[1-5]	Supports chaos engineering for AI agents?	[Documentation on test harnesses]
Escalation Paths for Critical Failures	[1-5]	Defines clear escalation to [CRITICAL INCIDENT TEAM].	[Runbook or escalation policy]

3.6 PERFORMANCE & SCALABILITY

(Ensures the tool can handle production workloads.)

Criteria	Scoring (1-5)	Notes	Evidence Required
Latency for High-Volume Logs	[1-5]	Handles [LOG VOLUME] with <[LATENCY THRESHOLD] ms per query.	[Benchmark report or SLA]
Agent Concurrency Support	[1-5]	Supports [MAX CONCURRENT AGENTS] without degradation.	[Scalability whitepaper]
Data Retention Flexibility	[1-5]	Adjusts retention per [DATA CLASSIFICATION] (e.g., PII vs. operational).	[Retention policy docs]
Multi-Cloud/On-Prem Support	[1-5]	Deployable in [DEPLOYMENT ENVIRONMENT] (e.g., AWS, GCP, air-gapped).	[Deployment guides]
Vendor SLAs for Uptime	[1-5]	Guarantees [UPTIME SLA]% with downtime compensation.	[SLA document]

3.7 SECURITY & PRIVACY

(Protects sensitive data and prevents misuse.)

Criteria	Scoring (1-5)	Notes	Evidence Required
Data Encryption (At Rest & In Transit)	[1-5]	Uses [ENCRYPTION STANDARD] (e.g., AES-256, TLS 1.3).	[Security whitepaper or compliance certs]
Data Minimization	[1-5]	Only collects [NECESSARY DATA FIELDS] for observability.	[Privacy impact assessment or data flow diagram]
Sensitive Data Masking	[1-5]	Redacts [SENSITIVE DATA PATTERNS] (e.g., PII, credentials) in logs.	[Demo of masking or API docs]
Third-Party Security Audits			

KESTRELATTICE

The governed agent mesh for production AI.

kestrelattice.com