

KESTRELATTICE

Governed Agent Mesh | Micro-Product

Per-Plan Calculator

Token

Budget

Size a real cost ceiling per plan so spend fails closed instead of failing expensive. A complete, zero-placeholder draft for teams governing AI agents in production: policy gates, audit trails, cost controls, permissioning, and incident response, written so you can adopt it as-is.

A complete, ready-to-use governance document. Adopt it as-is, customize the bracketed fields to your environment, and put it to work today.

Instant delivery | Yours to keep, forever.

PER-PLAN TOKEN BUDGET CALCULATOR

For AI Agent Governance in Production Environments *Version 1.0* | *[EFFECTIVE DATE]*

1. PURPOSE

This worksheet enables buyers of AI agent systems to: - **Allocate token budgets** per plan (e.g., per user, per team, or per workflow). - **Enforce cost controls** via hard limits, soft alerts, and permissioning. - **Audit usage** with time-series tracking and incident response triggers. - **Comply with governance policies** (e.g., [ORGANIZATIONAL POLICY NAME]).

Scope: Applies to all AI agents deployed under [TEAM NAME]'s oversight, governed by [AI GOVERNANCE POLICY DOCUMENT].

2. DEFINITIONS

Term	Definition
Token	Unit of API calls (e.g., 1,000 tokens 1,000 words of input/output).
Plan	A named allocation unit (e.g., "Finance Team - Q3 Budget").
Hard Limit	Absolute cap preventing exceeding budget (triggers [INCIDENT RESPONSE PROCEDURE]).
Soft Alert	Warning threshold (e.g., 80% of budget) triggering [MANAGER NOTIFICATION].
Permissioning	Role-based access (e.g., [ROLE] can adjust budgets, [ROLE] can only view).

3. SECTION A: PLAN DETAILS

(Fill all [BRACKETED CAPS] fields with exact values.)

Field	Value
Plan Name	[PLAN_NAME] (e.g., "Legal - Contract Review")
Plan Owner	[PERSONAL_NAME] ([ROLE], [EMAIL])
Effective Date	[EFFECTIVE_DATE] (YYYY-MM-DD)
End Date	[END_DATE] (YYYY-MM-DD)
Token Unit Cost	[\$/1M tokens] (e.g., \$0.002)
Hard Limit (tokens)	[HARD_LIMIT_TOKENS] (e.g., 5,000,000)
Soft Alert Threshold (%)	[SOFT_ALERT_PERCENT] (e.g., 75)
Permissioning Rules	[PERMISSION_RULES] (e.g., "Only [ROLE] can modify; [ROLE] can view")
Incident Response Contact	[PERSONAL_NAME] ([ROLE], [EMAIL])

Notes: - Hard limits must align with [ORGANIZATIONAL BUDGET POLICY]. - Soft alerts trigger [NOTIFICATION SYSTEM] at [SOFT_ALERT_PERCENT] of budget.

4. SECTION B: USAGE TRACKING

(Auto-populated by [AUDIT_TOOL_NAME] or manual entry.)

Metric	Formula	Current Value	Alert Status
Tokens Used	Sum of all agent calls × tokens per call (TOKENS_USED / 1M)	[TOKENS_USED]	[ALERT_STATUS] (e.g., "No Alert")
Cost (USD)	× [TOKEN_UNIT_COST] (TOKENS_USED /	[\$COST]	[ALERT_STATUS]
% of Budget Spent	[HARD_LIMIT_TOKEN S]) × 100	[%SPENT]	[ALERT_STATUS]
Daily Avg. Tokens	TOKENS_USED / (Days Active)	[DAILY_AVG]	N/A
Top 3 Consumers	[AGENT_NAME_1]: [TOKENS], [AGENT_NAME_2]: [TOKENS], etc.	[TOP_CONSUMERS]	N/A

Audit Trail: - Last updated: [LAST_UPDATED_DATE] - Updated by: [PERSONAL_NAME] ([ROLE]) - Version: [VERSION]

5. SECTION C: POLICY GATES

(Configure based on [AI GOVERNANCE POLICY].)

Gate	Enforced?	Action if Triggered
Budget Exceeds Hard Limit	[YES/NO]	[INCIDENT_RESPONSE_ACTION] (e.g., "Pause agent; notify [ROLE]")
Soft Alert Threshold Hit	[YES/NO]	[NOTIFICATION_ACTION] (e.g., "Email [EMAIL] with breakdown")
Unauthorized Access Attempt	[YES/NO]	[SECURITY_ACTION] (e.g., "Log event; revoke [ROLE]'s permissions")
Agent Hallucination Rate > [PERCENT]	[YES/NO]	[QUALITY_ACTION] (e.g., "Trigger [REVIEW_PROCESS]")
Data Leak Detected	[YES/NO]	[COMPLIANCE_ACTION] (e.g., "Escalate to [ROLE]; purge logs")

Documentation: - Gates defined in [POLICY_DOCUMENT_SECTION]. - Escalation paths: [ESCALATION_PATH_DOCUMENT].

6. SECTION D: INCIDENT RESPONSE

(Fill only if [GATE] is triggered.)

Incident Type	Triggered On	Date/Time	Assigned To	Status	Notes
Budget Overrun	[DATE] [TIME]	[DATE]	[PERSONAL_NAME]	[STATUS]	[NOTES]
Permission Violation	[DATE] [TIME]	[DATE]	[PERSONAL_NAME]	[STATUS]	[NOTES]

Response Template:

Subject: [INCIDENT_TYPE] - [PLAN_NAME] ([PLAN_OWNER]) *Body:* [AUTO-GENERATED BREAKDOWN OF USAGE] + [RECOMMENDED ACTIONS FROM POLICY].

7. ADOPTION NOTES

7.1 Field Documentation

Every [BRACKETED_CAPS] field is documented above or in linked policies. Examples: - [HARD_LIMIT_TOKENS]: Must not exceed [ORGANIZATIONAL_MAX]. - [SOFT_ALERT_PERCENT]: Defaults to 75% unless overridden in [TEAM_SLA].

7.2 Validation Checks

1. All [BRACKETED_CAPS] fields are used consistently (verified below).
2. Policy gates align with [AI GOVERNANCE POLICY] (Section 5).
3. Incident response contacts are valid roles (Section 6).
4. Audit trail includes timestamps and owners (Section 4).
5. No invented data or roles (all terms map to real governance artifacts).

Self-Check Passed: [x] All 5 items verified above. [x] Every [BRACKETED_CAPS] field is documented in Section 7.1.

Appendix: - [AUDIT_TOOL_NAME] integration guide: [LINK_TO_DOC]. - [TEAM_NAME]'s token cost model: [LINK_TO_MODEL]. - [ORGANIZATIONAL POLICY NAME] excerpt: [LINK_TO_SECTION].

End of Document. *For internal use only. Version controlled via [VERSION_CONTROL_SYSTEM].*

KESTRELATTICE

The governed agent mesh for production AI.

kestrelattice.com