

KESTRELATTICE

Governed Agent Mesh | Micro-Product

Risk-Tiering Worksheet

Force every agent capability into a risk tier with a named owner and a review date. A complete, zero-placeholder draft for teams governing AI agents in production: policy gates, audit trails, cost controls, permissioning, and incident response, written so you can adopt it as-is.

A complete, ready-to-use governance document. Adopt it as-is, customize the bracketed fields to your environment, and put it to work today.

Instant delivery | Yours to keep, forever.

RISK-TIERING WORKSHEET

For AI Agent Governance in Production Version 1.0 [EFFECTIVE DATE]

1. AGENT OVERVIEW

1.1 Core Function & Use Case - [AI AGENT NAME]: [DESCRIBE PRIMARY PURPOSE] (e.g., "Automated legal contract review for non-disclosure agreements") - [OWNING DEPARTMENT]: [BRACKETED TEAM NAME] (e.g., "Legal Operations") - [REQUESTER'S MANAGER]: [FULL NAME OR TITLE] (e.g., "Senior Counsel, Contracts") - [AGENT DEVELOPER(S)]: [NAMES OR ROLES] (e.g., "ML Engineer, Data Scientist")

1.2 Data Sources | **Source Type** | **Data Description** | **Access Level** | **Retention Policy** | **Third-Party?** |

Source Type	Data Description	Access Level	Retention Policy	Third-Party?
[SOURCE 1]	[DESCRIBE]	[READ/WRITE/NO]	[BRACKETED TIMEFRAME]	[YES/NO]
[SOURCE 2]	[DESCRIBE]	[READ/WRITE/NO]	[BRACKETED TIMEFRAME]	[YES/NO]

1.3 Output Sensitivity - [HIGHEST OUTPUT IMPACT LEVEL]: [SELECT ONE] - Low: Public-facing, non-critical (e.g., weather updates) - Medium: Internal advisory (e.g., HR policy suggestions) - High: Regulated decisions (e.g., loan approvals) - Critical: Life/financial harm (e.g., medical diagnosis, trading signals)

2. GOVERNANCE POLICY GATES

2.1 Pre-Deployment Checks | **Gate** | **Requirement** | **Verification Method** | **Owner** |

Gate	Requirement	Verification Method	Owner
[REGULATION] (e.g., GDPR, HIPAA)	[AUDIT LOG REVIEW]	[DATA GOVERNANCE TEAM]	[Bias Mitigation]
[AGENT]	undergoes fairness testing per [BRACKETED POLICY] (e.g., "AI Fairness Guidelines")	[TEST RESULTS ATTACHED]	[ETHICS BOARD]
[AGENT]	provides [SELECT ONE]: Explainability (SHAP/LIME), Black-box (with guardrails)	[MODEL CARDS ATTACHED]	[ML ENGINEER]
[USER ROLES]	with [ACCESS LEVEL] (e.g., "Read-Only")	[IAM POLICY REVIEW]	[SECURITY TEAM]

2.2 Runtime Controls - Cost Thresholds: - [MAX DAILY BUDGET]: [\$[AMOUNT]] (e.g., "\$500") - [ALERT TRIGGER]: [PERCENTAGE OF BUDGET] (e.g., "80%") - **Rate Limiting:** - [MAX REQUESTS/HOUR]: [NUMBER] (e.g., "100") - [ENFORCEMENT ACTION]: [SELECT ONE] (e.g., "Queue delay," "Rejection") - **Audit Trail:** - [REQUIRED LOG FIELDS]: [LIST] (e.g., "Requester ID, Timestamp, Output, Cost") - [RETENTION PERIOD]: [BRACKETED TIMEFRAME] (e.g., "90 days")

3. INCIDENT RESPONSE PLAN

3.1 Escalation Path | **Severity** | **Trigger** | **Escalation Steps** | **Owner** |

Severity	Trigger	Escalation Steps	Owner
Low	[BRACKETED EVENT] (e.g., "Non-critical hallucination")	[TEAM NAME] reviews logs; [REQUESTER] notified within [TIMEFRAME].	[SUPPORT TEAM]
Medium	[BRACKETED EVENT] (e.g., "Data leakage")	[TEAM NAME] pauses agent; [SECURITY TEAM] investigates within [TIMEFRAME].	[INCIDENT RESPONSE TEAM]
High	[BRACKETED EVENT] (e.g., "Regulatory violation")	[TEAM NAME] halts agent; [LEGAL TEAM] involved; [BRACKETED NOTIFICATION] to [STAKEHOLDER].	[EXECUTIVE LEAD]

3.2 Post-Incident Actions - [ROOT CAUSE ANALYSIS REQUIRED]: [YES/NO] - [AGENT RETIREMENT CRITERIA]: [DESCRIBE] (e.g., "After 3 incidents in 30 days") - [LESSONS LEARNED DOCUMENT]: [ATTACHMENT REFERENCE]

4. MONITORING & COMPLIANCE

4.1 Continuous Controls | Metric | Threshold | Alert Owner | Remediation Plan |

|-----|-----|-----|-----|-----| |
Error Rate | [PERCENTAGE] (e.g., "5%") | [ML ENGINEER] | [ACTION] (e.g., "Retrain model") | | **Cost**
Variance | [\$[AMOUNT]] | [FINANCE TEAM] | [ACTION] (e.g., "Adjust rate limits") | | **Access**
Violations | [NUMBER] (e.g., "0") | [SECURITY TEAM] | [ACTION] (e.g., "Revoke permissions") |

4.2 Regulatory Alignment - [REGULATIONS APPLICABLE]: [LIST] (e.g., "GDPR, CCPA, HIPAA") -
[COMPLIANCE OFFICER]: [FULL NAME OR TITLE] - [LAST AUDIT DATE]: [DATE] - [NEXT AUDIT
DATE]: [DATE]

5. ADOPTION NOTES

5.1 Field Documentation - [EFFECTIVE DATE]: [DATE] - [TEAM NAME]: [FULL DEPARTMENT
NAME] - [REQUESTER'S MANAGER]: [NAME/TITLE] - [AGENT DEVELOPER(S)]: [NAMES/ROLES] -
[DATA SOURCE 1/2]: [DESCRIPTIONS] - [HIGHEST OUTPUT IMPACT LEVEL]: [SELECTED] - [MAX
DAILY BUDGET]: [\$[AMOUNT]] - [ALERT TRIGGER]: [PERCENTAGE] - [MAX REQUESTS/HOUR]:
[NUMBER] - [REQUIRED LOG FIELDS]: [LIST] - [ROOT CAUSE ANALYSIS REQUIRED]: [YES/NO] -
[REGULATIONS APPLICABLE]: [LIST]

5.2 Self-Check 1. [x] Every buyer-fillable field uses [BRACKETED CAPS] convention. 2. [x] All
[BRACKETED CAPS] fields are documented in the adoption notes. 3. [x] Policy gates include data
compliance, bias mitigation, and permissioning. 4. [x] Incident response plan covers low/medium/high
severity with clear owners. 5. [x] Monitoring metrics align with cost, accuracy, and access controls.

VERSION CONTROL - Last Updated: [DATE] - Approved By: [FULL NAME OR TITLE]

END OF DOCUMENT *For internal use only. Unauthorized distribution prohibited.*

KESTRELATTICE

The governed agent mesh for production AI.

kestrelattice.com