

KESTRELATTICE

Governed Agent Mesh | Micro-Product

2026 09 29 Boardroom Briefing Pack.Md

A complete, ready-to-use governance document. Adopt it as-is, customize the bracketed fields to your environment, and put it to work today.

Instant delivery | Yours to keep, forever.

PRODUCT DRAFT -- The Boardroom Briefing Pack: AI Agent Governance for Executives & Boards

Status: draft -- never published, listed, uploaded, or sold. Human review and approval required before anything goes live.

Working title

The Boardroom Briefing Pack -- AI Agent Governance for Executives & Boards

One-line pitch

Board-ready briefings, one-pagers, and plain-English scripts so your security leader walks into the boardroom with answers about AI agents instead of anxiety.

Target buyer

The security/engineering leader or CISO who must explain their AI agent program's risks and controls to non-technical executives and the board -- credibly, in under twenty minutes.

Table of contents (draft product, ~15 pages)

1. **Read me first: how to use this pack** -- which piece to pull for a board slot, an audit prep call, or a new-exec onboarding; the 5-minute, 10-minute, and 20-minute tracks.
2. **The 10-minute verbal briefing script** -- a read-aloud script in plain English: what AI agents are, the three real risk categories, the three risk tiers, and the three questions the board should ask at the end. (*Sample section included below.*)
3. **The one-page executive risk summary** -- a filled example one-pager (based on the Studio Edition's risk-tier classification) plus a blank version: program scope, tiered actions, spend posture, audit posture, open items.
4. **The quarterly governance report template** -- one template, one filled example: what changed this quarter (new agents, new integrations, re-tiered actions), approval volume, blocked actions and why, spend vs. ceiling, audit-log health, next quarter's plan.
5. **"The five questions the board will ask" cheat sheet** -- who approves what, what our per-workload spend ceiling is, how we'd prove to an auditor what an agent did last Tuesday, what happens when a limit is hit, who owns re-tiering when integrations change -- each answered in two or three plain-English sentences.
6. **Plain-English glossary** -- agent, API, approval gate, grant, audit row, fail-closed, bounded plan -- one sentence each, zero jargon.
7. **Appendix: where each control comes from** -- maps every briefing claim to the Kestrelattice 8-layer governance stack and to the Studio Edition policy templates, so a technical skeptic on the board can trace any statement to a control.

Source-material grounding

Every control referenced in the pack traces to real Kestrelattice material: - Risk tiers (Hard-gated / Bounded auto / Draft-only) -- free playbook §4 and Studio Edition Part 4 (risk-tier classification worksheet with 8 pre-filled rows). - Hard-gated action list (money, contracts, external outreach, credentials, permission changes) -- Studio Edition Template 3. - Spend-ceiling patterns and the \$700-in-72-hours incident -- free playbook §5 and the evidence report (Reddit r/IndieDev). - Tamper-evident audit record requirement ("it probably didn't run" is not an audit trail) -- free playbook §6 and Studio Edition Template 6 / audit-row schema. - The 40%-by-2026 enterprise agent adoption

forecast -- Gartner, cited in the evidence report. No statistics, testimonials, or customer stories are invented anywhere in the pack.

SAMPLE SECTION -- Section 2: The 10-Minute Verbal Briefing Script

(Full section, as it would ship. Generated via the fleet dispatcher from the grounding facts above; statistics corrected against the evidence report: \$700+ burn, no invented per-workload ceiling.)

The 10Minute Verbal Briefing Script

Good morning, everyone. I'll give you a quick, plainEnglish overview of our AI agent program, the risks we face, and the safeguards we have in place. I'll keep it to about ten minutes, so feel free to ask questions afterward.

What are AI agents? An AI agent is a software program that can read data, make decisions, and perform actions on its own--like a digital employee. It can talk to other programs, send emails, or even move money, all without a person pressing a button. Think of it as a very smart, autonomous tool that can help us get work done faster.

Three real risk categories

Spending real money on API calls An AI agent can use a paid service (an "API") to do things like look up information or process payments. If the agent gets stuck in a loop, it can keep calling the service and burn up money. A public report documented a developer losing \$700+ in 72 hours because a runaway agent kept making calls. *Risk:* Unexpected costs.

Taking highconsequence actions AI agents can do things such as moving money, signing contracts, or sending messages to customers. Every time an agent performs an action that could materially affect the company, we must record a human approval before it happens. *Risk:* Unchecked actions that could harm the business.

Auditability Auditors need a tamperevident record--a clear, unalterable log--of what an agent did and why. A vague statement like "it probably did not run" is not enough. We must be able to show exactly what the agent did, when it did it, and what decision led to that action. *Risk:* Inability to prove compliance or investigate incidents.

Risk tiers

Hardgated In this tier, any highconsequence action--like moving money or signing a contract--must have a human signoff before the agent can act. There are no exceptions; the agent cannot bypass the gate. *Why it matters:* We keep full control over critical decisions.

Bounded auto Agents in this tier can run routine work automatically, but only within preapproved limits. We set a budget cap and a time limit. If the agent reaches either limit, it stops automatically. *Why it matters:* We allow efficiency while preventing runaway spending or time usage.

Draftonly Here the agent prepares drafts--emails, proposals, or reports--but never sends them, moves money, or contacts anyone. A human reviews the draft and decides whether to act. *Why it matters:* We use AI to accelerate creation but keep final authority with people.

Key points for the board

1. **Which actions require human approval?** Any action that could affect money, legal status, or external communication must be preapproved by a designated person.
 2. **What is our perworkload spend ceiling?** We have set a maximum spend of [YOUR PERWORKLOAD SPEND CEILING] for boundedauto tasks.
 3. **How would we prove to an auditor what an agent did last Tuesday?** We maintain an immutable, tamperevident log that records every API call, decision, and action, along with timestamps and the approving human's identity.
-

Gartner Insight Gartner projects that by the end of 2026, 40% of enterprise applications will embed taskspecific AI agents. This highlights the importance of having solid governance now.

Thank you for your attention. I'm happy to answer any questions you have.

-- [Your Name], CISO

Suggested price: \$19

Rationale. \$19 keeps this pack in the same tier as the six existing \$19 Kestrelattice kits, so bundle upgrade math stays clean (a buyer of any kit adds this without sticker shock). It also matches the market anchor: the packs competitors charge \$29+/seat/mo for observability tooling sit far above any one-time \$19 template pack, so this prices as an impulse add-on rather than a platform decision.

Gumroad-ready product description (~100 words)

Your AI agents are doing things. Can you explain their risks -- in plain English -- to the people who govern your company?

The Boardroom Briefing Pack gives you everything you need to brief executives and the board on AI agent governance: a 10-minute read-aloud briefing script, a filled one-page executive risk summary, a quarterly governance report template, a cheat sheet answering the five questions the board will ask, a plain-English glossary, and an appendix tracing every claim to a concrete governance control.

No jargon, no slides to design, no consultant required. Fill in your company's details, walk in, and own the conversation. A Kestrelattice governance template pack -- 100% self-serve digital download.

Draft metadata

- Bot: kestrelattice-product-forge
- Run date: 2026-09-29
- Generated sample section via: fleet dispatcher (groq:openai/gpt-oss-20b), with statistics corrected against the evidence report
- Ledger entry: appended to ledger.jsonl with status "draft"

KESTRELATTICE

The governed agent mesh for production AI.

kestrelattice.com