

KESTRELATTICE

Governed Agent Mesh | Micro-Product

Emergency Override Exception Template

The one form for breaking the rules in an emergency -- with automatic expiry and review. A complete, zero-placeholder draft for teams governing AI agents in production: policy gates, audit trails, cost controls, permissioning, and incident response, written so you can adopt it as-is.

A complete, ready-to-use governance document. Adopt it as-is, customize the bracketed fields to your environment, and put it to work today.

Instant delivery | Yours to keep, forever.

EMERGENCY OVERRIDE EXCEPTION TEMPLATE

For AI Agents in Production Environments Version 1.2 | Governance Compliance & Risk Mitigation

1. DOCUMENT OVERVIEW

This template formalizes an **exception request** for AI agent behavior that deviates from predefined guardrails, cost controls, or operational constraints. It ensures accountability, auditability, and alignment with **enterprise risk policies** while maintaining compliance with **GDPR, CCPA, or sector-specific regulations** (e.g., HIPAA, PCI-DSS).

Purpose: - Justify deviations from standard agent behavior. - Document escalation paths for high-risk overrides. - Maintain a traceable audit trail for compliance and incident response.

Scope: Applies to **all AI agents deployed in [BRACKETED CAPS]** environments, including but not limited to: - [PRODUCTION ENVIRONMENT NAME] - [CLOUD PLATFORM NAME] (e.g., AWS, Azure, GCP) - [ON-PREMISES/EDGE DEPLOYMENT LOCATION]

2. REQUEST DETAILS

2.1 Core Exception Request

Field	Definition	Value
[EXCEPTION REQUEST ID]	Unique identifier (auto-generated or manual)	
[REQUESTER NAME]	Full name of the individual submitting the request	
[REQUESTER ROLE]	Job title (e.g., "ML Engineer," "Compliance Officer")	
[REQUESTER'S MANAGER]	Name of the requester's direct supervisor	
[EFFECTIVE DATE]	Date/time the override must activate (YYYY-MM-DD HH:MM)	
[EXPIRATION DATE]	Date/time the override must terminate (must be [MAX OVERRIDE DURATION])	
[OVERRIDE PRIORITY]	Severity level: Critical (P1), High (P2), Medium (P3), Low (P4)	

Field	Definition	Value
[JUSTIFICATION]	Detailed rationale (max 500 words) explaining why standard behavior is insufficient. Must include: - Business impact (e.g., "Critical system outage during [EVENT NAME]"). - Risk assessment (e.g., "Override risk score: [CALCULATED RISK SCORE/100]"). - Temporary nature (e.g., "This is a one-time exception for [SPECIFIC USE CASE]").	

2.2 Agent-Specific Parameters

Field	Definition	Value
[AGENT NAME]	Exact name of the AI agent (e.g., "Fraud Detection Model v3.2")	
[AGENT VERSION]	Deployment version (e.g., "2024-03-15")	
[OVERRIDE TYPE]	Category of deviation: - Behavioral (e.g., "Bypass safety prompt") - Cost (e.g., "Increase API calls beyond [THRESHOLD]") - Data (e.g., "Process PII without encryption") - Output (e.g., "Generate unredacted responses")	
[SPECIFIC OVERRIDE ACTION]	Precise description of the deviation (e.g., "Disable rate_limit for 48h during peak traffic").	
[FALLBACK MECHANISM]	Plan if the override fails or causes harm (e.g., "Automatically revert to [FALLBACK AGENT]").	

2.3 Risk Mitigation & Controls

Field	Definition	Value
[RISK ASSESSMENT]	Structured risk matrix (use template below): Low/Medium/High	Likelihood Low/Medium/High

Field	Definition	Value
[COST ESTIMATE]	Monetary impact of the override (e.g., "\$12,450 in excess API costs"). Approval required if exceeding [MAX ALLOWED COST THRESHOLD].	
[AUDIT TRAIL REQUIREMENT]	Data to log: - [BRACKETED CAPS] (e.g., "All user IDs affected") - [BRACKETED CAPS] (e.g., "Timestamp of override activation") - [BRACKETED CAPS] (e.g., "Reason for termination").	
[PERMISSIONING]	Roles granted temporary access: - [ROLE 1] (e.g., "Data Scientist") - [ROLE 2] (e.g., "Compliance Officer") Expiration: [EXPIRATION DATE].	

3. APPROVAL WORKFLOW

3.1 Escalation Path

Approval Level	Role	Required Actions	Decision Timeframe
Tier 1	[REQUESTER'S MANAGER]	Review justification and risk matrix. Approve or escalate.	[BRACKETED CAPS] (e.g., "4 hours")
Tier 2	[TEAM LEAD, GOVERNANCE] AI	Verify technical feasibility and cost controls.	[BRACKETED CAPS] (e.g., "12 hours")
Tier 3	[CHIEF DATA OFFICER / CISO]	Final approval for high-risk overrides (e.g., data privacy violations). Sign-off for	[BRACKETED CAPS] (e.g., "24 hours")
Tier 4	[LEGAL/COMPLIANCE DEPT.]	regulatory-sensitive overrides (e.g., HIPAA-covered data).	[BRACKETED CAPS] (e.g., "48 hours")

3.2 Conditional Approvals

- **Cost Threshold:** If [COST ESTIMATE] > [MAX ALLOWED COST THRESHOLD], Tier 3 approval mandatory.
- **Data Sensitivity:** If override involves [BRACKETED CAPS] (e.g., "PII, PHI, or PCI data"), Tier 4 approval mandatory.

- **Critical Systems:** If agent affects [BRACKETED CAPS] (e.g., "Payment processing, patient records"). **Tier 3 + Tier 4 co-sign required.**

4. INCIDENT RESPONSE PLAN

Scenario	Detection Method	Escalation Path	Containment Steps	Post-Mortem Required?
Override silently fails	[BRACKETED CAPS] (e.g., "Agent logs")	[BRACKETED CAPS] (e.g., "PagerDuty")	[BRACKETED CAPS] (e.g., "Revert fallback")	[X]/[]
Cost budget exceeds	[BRACKETED CAPS] (e.g., "FinOps alert")	[BRACKETED CAPS]	[BRACKETED CAPS] (e.g., "Terminate override")	[X]/[]
Data leakage detected	[BRACKETED CAPS] (e.g., "SIEM trigger")	[BRACKETED CAPS]	[BRACKETED CAPS] (e.g., "Isolate affected data")	[X]
Agent generates harmful output	[BRACKETED CAPS] (e.g., "Safety monitor")	[BRACKETED CAPS]	[BRACKETED CAPS] (e.g., "Blacklist user input")	[X]

5. ADOPTION NOTES

5.1 Field Documentation

Field	Purpose	Who Fills It	Example Value
[EXCEPTION REQUEST ID]	Unique tracking for audit logs.	System (auto-generated)	EO-2024-0542
[REQUESTER'S MANAGER]	Ensures oversight of high-risk deviations.	Requester	"Dr. Emily Chen, VP of AI Ethics"
[MAX OVERRIDE DURATION]	Prevents indefinite exceptions.	Governance Policy Owner	72 hours
[THRESHOLD]	Cost/access limits to avoid abuse.	Finance/Compliance Team	\$5,000
[FALLBACK MECHANISM]	Minimizes downtime if override fails.	DevOps/SRE Team	"Revert to agent_v1.1"

5.2 Compliance Checklist

- [] [EXCEPTION REQUEST ID] is auto-generated or manually assigned.
- [] [JUSTIFICATION] includes business impact, risk score, and temporary justification.
- [] [RISK ASSESSMENT] table is fully populated with likelihood/impact/mitigation.

- [] Every [BRACKETED CAPS] field is documented in **Section 5.1**.
- [] All buyer-fillable fields use [BRACKETED CAPS] convention (verified: [x]).

6. APPENDICES

6.1 Risk Matrix Template

Likelihood	Impact	Risk Level	Mitigation Actions	Owner
Low	Low	Acceptable	Monitor; no action required.	[BRACKETED CAPS]
Low	High	Moderate	Implement [BRACKETED CAPS].	[BRACKETED CAPS]
High	High	Critical	Require [BRACKETED CAPS] approval.	[BRACKETED CAPS]

Risk Score Calculation: (Likelihood Score: 1-3) × (Impact Score: 1-3) = Risk Level (1-9) - **1

KESTRELATTICE

The governed agent mesh for production AI.

kestrelattice.com